

FinCatch  
Coding task – hard  
August 9, 2024

## Context

In the rapidly evolving landscape of financial technology, artificial intelligence plays a crucial role in forecasting, decision-making, and risk management. This assessment is designed to evaluate your ability to design, implement, and critically analyze AI systems in a financial context. You have 48 hours to complete this task, allowing you to showcase both your technical skills and your ability to manage a complex project within a realistic timeframe.

## Problem statement

You are tasked with developing an AI system for a quantitative trading firm. The system consists of three main components: a Large Language Model (LLM) for market analysis, an Agent for decision-making, and an Environment that simulates market conditions. Your goal is to implement these components, analyze their performance, and propose improvements.

### Q1 MLP Network for LLM Simulation (Market Analyzer)

Implement a Multi-Layer Perceptron (MLP) to simulate an LLM that analyzes market conditions.

- Input: A number between 1 and 100, representing various market indicators.
- Output: A number  $k \in [1, 10000]$ , representing the LLM's market analysis.
- Special cases:
  - When input is 1-9, output should be 1111-9999 respectively.
  - For inputs 10-100, output should be any other value in the range.

### Q2 Agent Policy Network (Decision Maker)

Design a network that takes the LLM output and decides on a trading action  $\alpha \in [1, 1000]$ . Implement the following reward structure:

- If LLM output is XXXX and action  $\alpha$  contains:
  - One X: 10 points
  - Two X's: 20 points
  - Three X's: 100 points
  - All other cases: 0 points

For example:

- Input 1  $\rightarrow$  LLM  $\rightarrow$  state = 1111  $\rightarrow$  Agent  $\rightarrow$  Action = 111  $\rightarrow$  returns reward = 100;

- Input 2  $\rightarrow$  LLM  $\rightarrow$  state = 3333  $\rightarrow$  Agent  $\rightarrow$  Action = 313  $\rightarrow$  returns reward = 20.
- Input 3  $\rightarrow$  LLM  $\rightarrow$  state = 3333  $\rightarrow$  Agent  $\rightarrow$  Action = 124  $\rightarrow$  returns reward = 0.

### Q3 Environment simulation (Market simulator)

In real-world scenarios, you may not have a predefined reward function, and the Agent needs to interact with the environment multiple times before receiving a reward. Design an environment function, a reward function, and use reinforcement learning to complete the following test cases:

*Test Case 1:*

Input 1  $\rightarrow$  LLM  $\rightarrow$  state = 1111  $\rightarrow$  Agent  $\rightarrow$  Action = 111  $\rightarrow$  Environment: returns the number of matching digits  $\rightarrow$  State = 3  $\rightarrow$  LLM  $\rightarrow$  state = 3333  $\rightarrow$  Agent  $\rightarrow$  Action = 333  $\rightarrow$  Environment: returns the number of matching digits  $\rightarrow$  State = 3  $\rightarrow$  Done: state matches the previous result  $\rightarrow$  If state = 3, return reward =  $100/2$ , where 2 represents the number of interactions with the environment.

*Test Case 2:*

Input 1  $\rightarrow$  LLM  $\rightarrow$  state = 1111  $\rightarrow$  Agent  $\rightarrow$  Action = 199  $\rightarrow$  Environment: returns the number of matching digits  $\rightarrow$  State = 2  $\rightarrow$  LLM  $\rightarrow$  state = 2222  $\rightarrow$  Agent  $\rightarrow$  Action = 213  $\rightarrow$  Environment: returns the number of matching digits  $\rightarrow$  State = 1  $\rightarrow$  LLM  $\rightarrow$  state = 1111  $\rightarrow$  Agent  $\rightarrow$  Action = 111  $\rightarrow$  Environment: returns the number of matching digits  $\rightarrow$  State = 3  $\rightarrow$  LLM  $\rightarrow$  state = 3333  $\rightarrow$  Agent  $\rightarrow$  Action = 333  $\rightarrow$  Environment: returns the number of matching digits  $\rightarrow$  State = 3  $\rightarrow$  Done: state matches the previous result  $\rightarrow$  If state = 3, return reward =  $100/4$ , where 4 represents the number of interactions with the environment.

- If the interaction with the environment reaches 10 rounds without stopping, return reward = 0.
- If at the end, the state is not equal to 3, return reward = 0.

In shorts, the episode should end when:

- The state matches the previous state (successful prediction) 10 rounds of interaction are reached
- If final state is 3:  $100 / (\text{number of interactions})$
- Otherwise: 0

### Discussion on the question

If the policy network above is not "pre-trained", you might find it difficult to converge during training. Why do you think this happens? Can you suggest clever techniques/methods to efficiently train the Agent's Policy network? Provide analysis and code results.

### Future directions and open-ended exploration

What insights into AI agent design did you gain from the above tasks? Please use any open-source finance-related projects, dataset, and packages to make the agent more realistic. Provide your attempts and analysis.

## **Reminder**

Provide well-commented code for all implementations, and include a comprehensive report describing your approach, analysis, and findings.

There are no strictly right or wrong answers for the open-ended sections. We're interested in your thought process, creativity, and ability to tackle complex problems in AI and finance.

**Good luck!**